



# Final Demo

## Building Cyber Infrastructure for Researchers

### Mentors:

Abraham Matta and Ali Raza

### Team Members:

Tian Chen, Donovan Jones, Komal Kango, Jing Song and Kristi Perreault

# Project Overview



Create Infrastructure for Researchers from the Earth Science Department at BU that allows them to run submitted code on large data sets and display the results.

Targets end-users in the Earth and Science Department

Users do not include:


- Non- ecological Researchers
- Advanced users with complex requirements beyond the scope of the project.

# Vision & Goals



- Provide a web service with a simple user experience such that researchers can submit code.
- Develop reliable infrastructure on unreliable nodes using a Kubernetes Cluster
- Focus on Function as a Service with OpenWhisk as a proof of concept
- Provide a user interface that allows for comparisons between multiple models on the same data set along with comparisons of models using periodic data sets in order to determine model accuracy.

# System Diagram



# Compute with OpenWhisk



- Using Helm to enable Kubernetes cluster with OpenWhisk
- OpenWhisk API exposed externally with nginx
- Working with Chameleon/GENI to add/remove worker nodes on a Kubernetes cluster
- Complex Request/Response with OpenWhisk on a cluster
- Working with Plotly & MongoDB for displaying results

# Orchestration with Kubernetes



- Allows for a consistent layer over which anything can be deployed
- Focus is Function as a Service (FaaS) with OpenWhisk
- Install OpenWhisk on the cluster using Helm
- Deploy Kubernetes Cluster on the Mass Open Cloud

# Unreliable Nodes



- Utilize virtual machine in the MOC along with Chameleon and GENI as the unreliable nodes to build reliable infrastructure
- Capability for infrastructure to "loan out" these nodes to services or applications as needed\*
- Monitor the availability of these nodes, including up/down time and proximity to data stores\*

(\*) Indicates stretch goal

# OpenWhisk & Kubernetes Demo

```
ubuntu@kubernetes-cluster: ~$ sudo helm status owdev -n openwhisk
NAME: owdev
LAST DEPLOYED: Thu Apr 30 19:05:41 2020
NAMESPACE: openwhisk
STATUS: deployed
REVISION: 5
NOTES:
Apache OpenWhisk
Copyright 2016-2018 The Apache Software Foundation

This product includes software developed at
The Apache Software Foundation (http://www.apache.org/).

To configure your wsk cli to connect to it, set the apihost property
using the command below:

    $ wsk property set --apihost 172.17.0.2:31001

Your release is named owdev.

To learn more about the release, try:

    $ helm status owdev [--tls]
    $ helm get owdev [--tls]

Once the 'owdev-install-packages' Pod is in the Completed state, your OpenWhisk deployment is ready to be used.

Once the deployment is ready, you can verify it using:

    $ helm test owdev [--tls] --cleanup
```

```
ubuntu@kubernetes-cluster: ~$ sudo helm status owdev -n openwhisk
```



# User Interface



- User registration & login
- User hierarchy with system admin, project leads, project members
  - System Admin: manage all projects and users
  - Project Lead: Add/remove team members
- Request to join projects
- Code submission with codemirror editor
- Results visualization with plotly
- Ability to view previous computations

# Database Management




- User information stored and managed in MongoDB
- Store output of computation in MongoDB as a JSON object

# ecoforecast.bu.edu

- Currently under deployment

# User Interface Demo



# New feature compare to the previous project



Old project:

Web server written by **python script**

Disadvantages: difficult to test and debug; different to add new functions

---

Our project:

Web server written by **flask**

Advantages: Easy to test and debug; easy to extent new features

# New feature compare to the previous project



Old project:

Creating a thread to start OpenWhisk Command-line Interface to connect to the openwhisk

Disadvantages: Can't get the data directly from openwhisk; Do not know whether the request was failed

---

Our project:

Using OpenWhisk API

Advantages: Easy to test and debug; get information directly from openwhisk including fail information

# New feature compare to the previous project



Old project:

Deploying OpenWhisk on GENI or chameleon

Disadvantages: resource will expire and need to apply again, user can't access to openwhisk during this time

---

Our project:

Deploy OpenWhisk on GENI or chameleon, GENI and chameleon node managed by kubernetes

Advantages: we can add new node to the kubernetes and not affect the current system in the meantime

# What we learned this semester



- How to install and work with OpenWhisk
- Testing API calls with Postman
- Working with VMs and security on the MOC
- Standing up a Kubernetes node and cluster
  - Using Helm to install OpenWhisk
- Flask, HTML, and Python for the user interface
- Using Plotly & MongoDB for displaying results



# Semester Problems & Limitations



- Access and working with the MOC was difficult
  - Security issue
- Getting up to speed with how to use OpenWhisk took a while
- Working with Helm to create Kubernetes cluster took some time
- Ran out of space on the Kubernetes cluster when trying to run OpenWhisk commands
- COVID-19 repercussions led to many issues
  - Difficult to find time to meet & work together with time zones
  - Had to adjust project deliverables
  - Constant communication
  - Mental & physical health

# Semester Summary



- Created a Kubernetes cluster on MOC to facilitate Chameleon & GENI worker nodes
- Installed OpenWhisk on cluster to run functions for researchers
- Created MongoDB databases for storing user data and computation results
- Created a new, cleaner UI featuring:
  - User login & registration
  - Dashboard
  - User hierarchy (admin, project leads, project members)
  - Code submission and results

# Future Work



- Results comparison with previous computations
- More plotting options to view data in different ways
- More Kubernetes clusters, worker nodes spread geographically
- System Admin UI
  - View which nodes are up or down
  - See geographic distribution of data & nodes
- Monitoring Application
  - Monitor worker nodes
    - If expiring, get new resource from Geni or Chameleon and add worker
  - Interface that resources from the cloud and adds it to system as spot instances



**Thank you!**